

# Distributed Computing

## **A-01. Introduction to the Course**

# Welcome!

- This course will be shared between the
  - 1<sup>st</sup> year of DSE MSc
  - 1<sup>st</sup> year of SSE MSc
  - 2<sup>nd</sup> year of SSE MSc
- It's organized in a rather complex way, so we'll explain how it will work for everybody

# What Is a Distributed System?

- A collection of **autonomous computing elements (nodes)** that appear to its users as a **single coherent system** ([van Steen & Tanenbaum](#))
  - Colored text are links... Please follow them to know more!
- Distributed systems are **everywhere**: pretty much anything you see on the Internet is one...
- Also, a single multi-core computer can be seen as a distributed system; distributed system techniques can be and are applied even there ([Bauman et al.](#))

# Do I Need A Distributed System?

- **Availability:** if one computer (or 10) break down, my website/DB/fancy Ethereum app will still work
- **Performance:** a single-machine implementation won't be able to handle the load/won't be fast enough
- **Decentralization:** I don't want my system to be controlled by a single entity
- And all these things are related in non-trivial ways...  
**We'll see!**

# Making a System Coherent

- We've seen that distributed systems are about making disparate machines **coherent**
- To be coherent, nodes need to **collaborate**
  - We need **synchronization** (there is no **global clock**)
  - We need to **manage group membership & authorizations**
  - We need to deal with **node failures**

# Distributed Systems: a Huge Topic

- Many courses go in **depth**: you have a (small enough) topic and they teach you everything about it
- Here, we'll go in **breadth**: we'll introduce several topics at a high level, and go deeper on a few
  - Idea: giving you starting points for learning on your own, including after university

# How We See the Course

- Security Software and Engineering
  - You'll be among the ones **designing** and **securing** these systems
  - You'll be the “mechanics” who will be tinkering with these systems
- Data Science and Engineering
  - You'll be **using** these systems
  - You'll be the “pilots” who need to know their systems to use them well



# Organization of the Course

- Part A (**everybody**):
  - Distributed systems in general
  - Lessons held by Matteo Dell'Amico
- Part B1 (**DSE**):
  - Big Data Engines
  - Lessons held by Giorgio Delzanno
- Part B2 (**SSE** 2<sup>nd</sup> year):
  - Blockchains and Distributed Ledgers
  - Lessons held by Marina Ribaudó



# Exams: SSE

- You'll have **assignments** based on a simulator you have to complete
  - Questions like “in this scenario, which are the best design choices”?
- You'll **tinker** with the simulation and write reports answering the question
- Giacomo Benedetti will follow your assignments and give a green light to proceed to the exam
- Oral exams, where you'll present the assignments and will be asked questions about all the program

# Exams: DSE

- You'll have **exercises** on Apache Spark to do during the year
- A written final examination

# A Word About Theses

- More important for 2<sup>nd</sup> year students
- But consider what you're enjoying
- If you find some of these topics interesting, come talk to us!

# Part A: Some Topics We'll Touch

- Making systems **consistent**
  - **Consensus** mechanisms
- What **queueing theory** tells us
  - Effects of sharing load between servers
- Handling data efficiently
  - Modeling systems with data **replication**
  - **Erasure coding**: the gifts of coding theory
- Introduction to decentralized systems
- Introduction to big data engines